

CMPSC174A:Section4

RelationalAlgebra, Datalog

Feb 3rd, 2021

Administrivia

- HW3 due **next Monday**, Feb. 8th @ 11:00pm

RA Operators

$$R1 \cap R2 = R1 - (R1 - R2)$$

$$R1 \cap R2 = \cap \text{Intersect}$$

$$R1 \bowtie R2$$

Standard:

U - Union

σ - Select

π - Project

ρ - Rename

Joins:

\bowtie - Nat. Join

\square - L.O. Join

\square - R.O. Join

\square - F.O. Join

\times - Cross Product

Extended:

δ - Duplicate Elim.

γ - Group/Agg. τ -
Sorting

Y Notation

Grouping and aggregation on group:

$\text{Y}_{\text{attr}_1, \dots, \text{attr}_k, \text{count/sum/max/min}(\text{attr}) \rightarrow \text{alias}}$

Aggregation on the entire table:

$\text{Y}_{\text{count/sum/max/min}(\text{attr}) \rightarrow \text{alias}}$

Query Plans

Select-Join-Project structure

Make this SQL query into RA

```
SELECT R.b, T.c, max(T.a) AS T_max FROM
  Table_R R, Table_T T
  WHERE R.b = T.b
GROUP BY R.b, T.c      HAVING max(T.a) >
99
```

Query Plans

Select-Join-Project structure

Make this SQL query into RA

```
SELECT R.b, T.c, max(T.a) AS T_max  
FROM Table_R R, Table_T T  
WHERE R.b = T.b GROUP BY R.b,  
T.c HAVING max(T.a) > 99
```

$$\pi_{R.b, T.c, T_max}(\sigma_{T_max > 99}(\gamma_{R.b, T.c, \max(T.a) \rightarrow T_max}(R \bowtie_{R.b=T.b} T)))$$

Datalog Terminology

Head - Body - Atom/Subgoal/Relational predicate

Base Relations (EDB) vs Derived Relations (IDB)

- Negation + Aggregate

Query Safety

Need a positive relational atom of every variable

What's wrong with this query?

Find all of Alice's children without children:

```
U(x) :- ParentChild("Alice", x), !ParentChild(x, y)
```


Query Safety

```
U(x) :- ParentChild("Alice",x), !ParentChild(x,y)
```

It is domain dependent! Unsafe!

Double negation to the rescue. Why does this work?

```
NonAns(x) :- ParentChild("Alice",x), ParentChild(x,y)
```

```
# All of Alice's children with children
```

```
U(x) :- ParentChild("Alice",x), !NonAns(x)
```

```
# All of Alice's children without children (safe!)
```

But we can do better...

Query Safety

But we can do better...

```
hasChild(x) :- ParentChild(x,_) # People with  
children  
U(x) :- ParentChild("Alice",x), !hasChild(x)  
# All of Alice's children without children (safe!)
```

Datalog with Recursion

Able to write complicated queries in a few lines

Graph analysis

Done with query once output does not change.

Stratified Datalog

Recursion might not work well with negation

E.g.

```
A(x) :- Table(x), !B(x)
```

```
B(x) :- Table(x), !A(x)
```

Solution: Don't negate or aggregate on an IDB predicate until it is defined
Stratified Datalog Query

Stratified Datalog

Only IDB predicates defined in strata 1, 2, ..., n may appear under ! or agg in stratum n+1

