

CMPSC 174A: Section 8

Transactions

March 3rd, 2021

Administrivia

HW5 Deadline: March 9th, 2021 11:00 PM PT

HW5 - Concurrent

```
from concurrent.futures import ProcessPoolExecutor
import os

def task():
    print("Executing our Task on Process: {}".format(os.getpid()))

def main():
    ## executor = ProcessPoolExecutor(max_workers=3)
    with ProcessPoolExecutor(max_workers=3) as executor:
        task1 = executor.submit(task)
        task2 = executor.submit(task)

main()
```

HW5 - apsw Transaction

this will be 3 separate transactions

```
db.cursor().execute("INSERT ...")
```

```
db.cursor().execute("INSERT ...")
```

```
db.cursor().execute("INSERT ...")
```

this will be one transaction

```
db.cursor().execute("BEGIN")      // <<<-
```

```
db.cursor().execute("INSERT ...")
```

```
db.cursor().execute("INSERT ...")
```


```
db.cursor().execute("INSERT ...")
```

```
db.cursor().execute("COMMIT")     //. <<<-
```

Serializability

Conflict serializable is stricter than serializable



I.e. Any schedule that is conflict serializable must be serializable.



Serializability

y

Checking for conflict serializability \rightarrow precedence graph and cycle checking



Serializability

S1: w1(Y); w2(Y); w1(X); w2(X); w3(X)


Are these serializable?
Conflict serializable?

S2: w1(Y); w2(Y); w2(X); w1(X); w3(X)

Serializability


S1: $w1(Y)$; $w2(Y)$; $w1(X)$; $w2(X)$; $w3(X)$

Conflict Serializable



S2: $w1(Y)$; $w2(Y)$; $w2(X)$; $w1(X)$; $w3(X)$

Serializable (but not conflict serializable)



2PL v.s. Strict 2PL

2P:

- In every transaction, all lock requests must precede all unlock requests
- Ensure Conflict Serializability
- Might not be able to recover (Dirty Read: Read on some write that gets rolled back)

Strict 2PL:

- Every lock each transaction is held until commit or abort
- Ensure Conflict Serializability
- Recoverable as each transaction does not affect others until commit/abort

2PL v.s. Strict 2PL

A New Problem: Non-recoverable Schedule

T1	T2
$L_1(A); L_1(B); \text{READ}(A)$ $A := A + 100$ $\text{WRITE}(A); U_1(A)$	$L_2(A); \text{READ}(A)$ $A := A * 2$ $\text{WRITE}(A);$ $L_2(B); \text{BLOCKED} \dots$
$\text{READ}(B)$ $B := B + 100$ $\text{WRITE}(B); U_1(B);$	$\dots \text{GRANTED}; \text{READ}(B)$ $B := B * 2$ $\text{WRITE}(B); U_2(A); U_2(B);$ Commit

Rollback

Isolation Level: Read Uncommitted

Write Locks? Strict 2PL

Read Locks? No (Immediate Read)

Problem: Dirty-Read

Reading uncommitted data that can be rolled back

Isolation Level: Read Uncommitted

Example

T1	T2
W(A)	
	R(A)
	W(B)
	Commit
R(B)	
Commit	

T2 is reading value of A updated by T1's write on A, but T1 has not committed yet.

The value of A read by T2 might not even be in the result.

Then T2's action can be influenced by such uncommitted data.

Isolation Level: Read Committed

Write Locks? Strict 2PL

Read Locks? Obtain before read, release after (No more dirty read)

Problem: Unrepeatable Read

The values of 2 reads on the same tuple can be different in the same transaction

Isolation Level: Read Committed

Example
Transaction:

$$S = \begin{bmatrix} T1 & T2 \\ R(A) & \\ & R(A) \\ & W(A) \\ & Com. \\ R(A) & \\ W(A) & \\ Com. & \end{bmatrix}$$

T1's first R(A) and T1's second R(A) might have different results.

Updated by T2's W(A).

Isolation Level: Repeatable Read

Write Locks? Strict 2PL

Read Locks? Strict 2PL (No more unrepeatable read) Same as Serializable if no insert or delete

Problem: Phantom Read

In the same transaction, some tuples appear sometimes and disappear other times

Isolation Level: Repeatable Read

Suppose there are two blue products, A1, A2:

Phantom Problem

T1

T2

```
SELECT *  
FROM Product  
WHERE color='blue'
```

```
INSERT INTO Product(name, color)  
VALUES ('A3','blue')
```

```
SELECT *  
FROM Product  
WHERE color='blue'
```


Isolation Level: Serializable

Not the same thing as Serializable schedule!!!

Write Locks: Strict 2PL

Read Locks: Strict 2PL

Predicate Lock/Table Lock (No Phantom)

Isolation Level: Serializable

Predicate Lock Example:

In Transaction T, we have a statement:

SELECT * FROM People WHERE age > 18;

In this case, the transaction will grab a predicate lock that prevent inserting and deleting tuples that can affect the predicate/statement.

In this case, the lock prevents inserting and deleting tuples with age > 18.

Isolation Level: Summary

Isolation Level	Dirty Reads	Nonrepeat- able Reads	Phantoms
Read Uncommitted	Allowed	Allowed	Allowed
Read Committed	Not Allowed	Allowed	Allowed
Repeatable Read	Not Allowed	Not Allowed	Allowed
Serializable	Not Allowed	Not Allowed	Not Allowed